

Tensors and Torch Operations

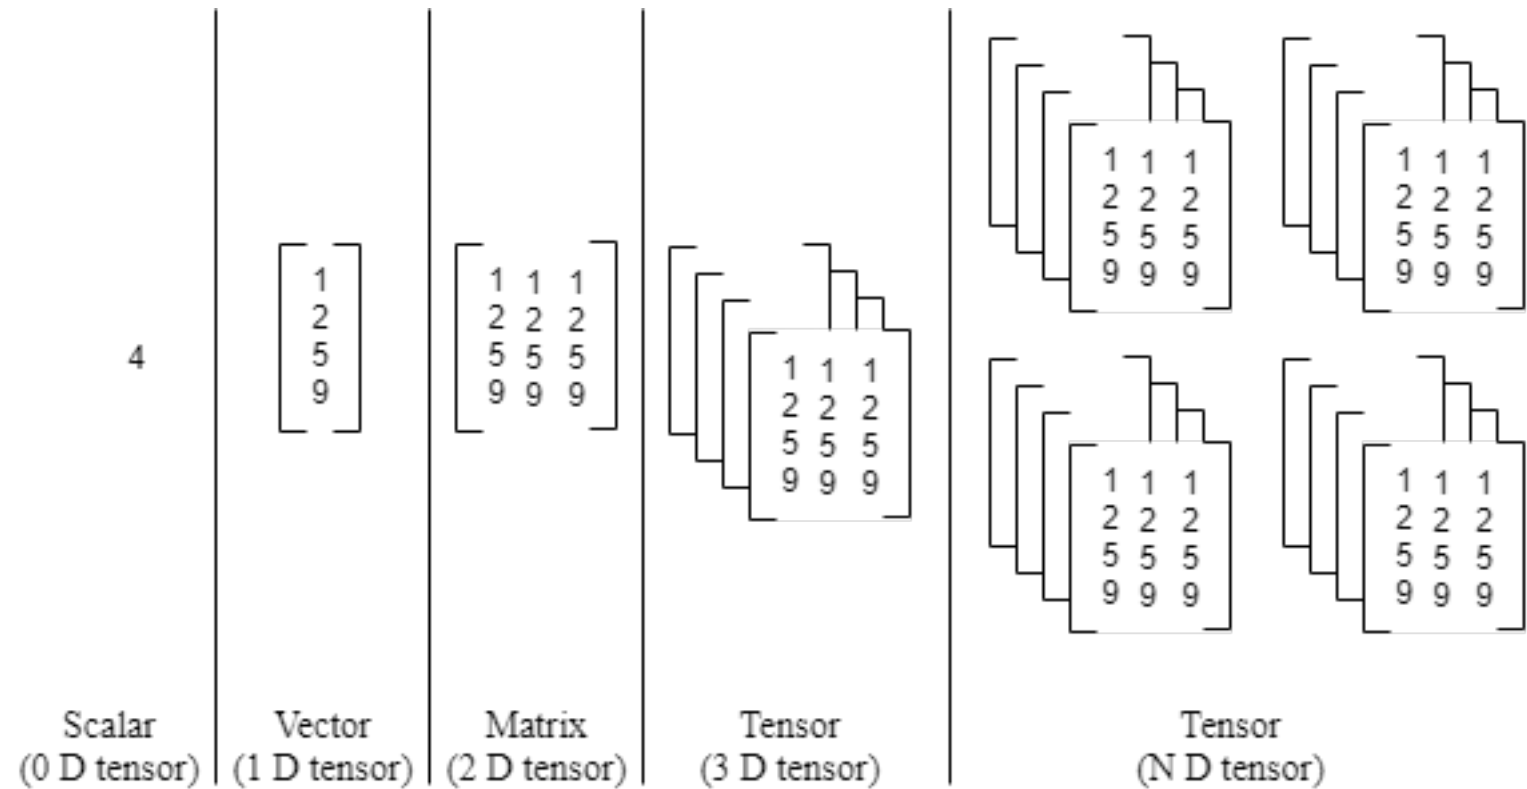
SYDE 599 Deep Learning F23

September 21, 2023



Tensors

- A tensor is a homogenous, n-dimensional array of numbers
 - 0-D: Scalar
 - 1-D: Vector
 - 2-D: Matrix
 - 3-D+: Tensor




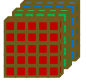
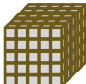
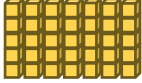


Tensor shape conventions



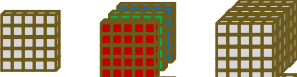
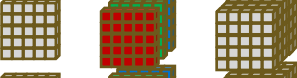
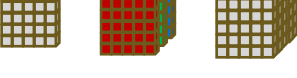
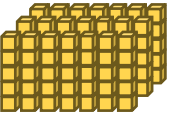
- In deep learning, we will use the following letters to represent certain dimensions
 - N for Number of examples, B for Batch size
 - L for Length of a sequence
 - W for Width, horizontal spatial dimension of an image
 - H for Height, vertical spatial dimension of an image
 - C for Channels (images or signals), D for feature Dimension (most other data)
- The batch dimension is always first / on the left
- The feature dimension is always last / on the right

Tensor shape conventions

Single data elements

- Scalar: $(1,)$ 
- Column vector: $(D,)$ or $(1, D)$ 
- Greyscale image: $(W, H, 1)$ 
- RGB image: $(W, H, 3)$ 
- Arbitrary image: (W, H, C) 
- Sequence of vectors: (L, D) 

Batches of data elements

- Row vector: $(B, 1)$ 
- Feature matrix: (B, D) 
- Greyscale images: $(B, W, H, 1)$ 
- RGB images: $(B, W, H, 3)$ 
- Arbitrary images: (B, W, H, C) 
- Sequences of vectors: (B, L, D) 

Torch dtypes

- We almost always want our tensor dtype to be `torch.float32`
 - Float operations default to `torch.float32`, it's also optimized for efficient compute on GPUs
 - Exception is arrays that are used for indexing, which use `torch.long` / `torch.int64`
 - By default, integers are `torch.long`, so conversion to `torch.float32` may be required

Tensor indexing



- Tensors can be indexed in the following ways to select elements:
 - Single integers (0, 1, 2, ..., D-1) or negative integers (-1, -2, ...) to select from end
 - Lists of integers (repeats allowed)
 - Slices, using colon notation (or slice objects)
 - Boolean masks the same size as the dimension(s)
 - Ellipsis to infer dimensions
- Size of resulting tensor after selection depends on method
 - Single integer: Dimension dropped
 - Multiple elements: Dimension kept, size is length of selector
 - Colon or ellipsis: Original dimension size(s) kept

Elementwise operations

- Easiest tensor operations to apply, they modify each element independently
- Examples: `torch.exp()`, `torch.relu()`, `torch.sigmoid()`

Broadcasting

- Typically, we would expect binary operators to work with two arrays when they have the exact same shape, so that the operation can be applied elementwise
- Broadcasting rules assume the shape of the right operand can be transformed to be compatible with the left operand




Broadcasting

- Broadcasting is compatible when:
 - Trailing (right-most) dimensions match left operand, or
 - A dimension has size 1
- Typical operations:
 - Addition: Adding a scalar to a tensor, adding a bias vector to each row
 - Multiplication: Multiplication by a scalar, scaling each column by its standard deviation
 - Exponentiation: Raise a tensor to power of a scalar, e.g. squaring
 - Comparison operations: Evaluate if all elements are less than a scalar, evaluate if elements equal zero

Matrix multiplication


- Also known as dot product or inner product
- Left operand (A) is at least 2-D tensor, right operand (B) is a 2-D tensor (matrix)
- Matmul is broadcastable if $\text{ndim}(A) > 2$
 - $(B, H, W, C) @ (C, D) = (B, H, W, D)$
- Matmul is a linear transformation on A from m-dimensional vector space to n-dimensional vector space



- Last dimension of A must be same size as second-last dimension of B
- C maintains same shape of A, but with last dimension the same as last dimension of B

Reduction operations

- Apply a function that reduces along one or more dimensions
- By default, reduces across entire tensor, but can specify dimension with keyword “dim” (note Numpy uses keyword “axis” instead)
- Reduction along a dim will drop that dim from the shape
- Most reduction operators use sum internally
- Examples:
 - `torch.sum()`, `torch.mean()`, `torch.std()`
 - `torch.softmax()`, loss functions like MSE, BCE, CE
 - `torch.min()`, `torch.max()`



Reshape operations

- Alter the shape / layout of the tensor without adding or modifying the elements
- Useful for making shapes compatible with broadcasting, especially if dims are lost after a reduction
- Examples:
 - `Tensor.reshape()`: General reshape, specify a tuple of the new shape
 - `Tensor.squeeze()`: Remove singleton dims
 - `Tensor.unsqueeze()`: Add a singleton dim
 - `Tensor.flatten()`: Unravel into a vector
 - `Tensor.permute()`: Permute order of dims
 - `Tensor.T`: Transpose, swap dims 0 and 1 (matrix only)