

Regularization and Optuna

SYDE 599 Deep Learning F23

October 26, 2023



Assignments

- If there are Assignment 1 questions, you can ask after class
- Self-enroll groups available for Assignment 2 (2-3 students) and Assignment 3 / Presentation (4-5 students)
- Assignment 2 should be easily completed with only 2 people, groups from Assignment 1 can split into two groups
- Completing A2 should help you feel ready to use PyTorch for the project

Optuna

- Not installed by default in Google Colab, install with `!pip install optuna`
- [Optuna documentation](#) -> Key features tutorials

Key Optuna Terminology

- Study
 - Hyperparameter optimization session, consisting of a set of trials
- Trial
 - Process of evaluating an objective function once
 - Provides interface to get suggested parameters based on optimization algorithm
 - Considerations
 - Range of values
 - Log scale (e.g. learning rate, L2 penalty)
 - Discrete or continuous

<code>suggest_categorical()</code>	Suggest a value for the categorical parameter.
<code>suggest_discrete_uniform(name, low, high, q)</code>	Suggest a value for the discrete parameter.
<code>suggest_float(name, low, high, *[, step, log])</code>	Suggest a value for the floating point parameter.
<code>suggest_int(name, low, high[, step, log])</code>	Suggest a value for the integer parameter.

Key Optuna Terminology

- Objective
 - Non-differentiable goal you want to optimize, e.g. maximize test accuracy on MNIST with a neural network
 - Single function, input is a “trial” and output is your metric
 - Suggest hyperparameters, construct model, perform training, and perform evaluation within the objective
- Sampler
 - Non-gradient based optimization method for (hyper)parameters
 - Consider reporting reasons for sampler choice in the report

Optuna Objective for Deep Learning

```
def objective(trial):  
    hps = suggest_hyperparameters(trial)  
    model = create_model(hps)  
    model = train_model(model, train_dataset, hps)  
    metric = evaluate_model(model, test_dataset)  
    return metric
```

Convolution in Torch

- Torch expects image tensors of shape (B, C, H, W)
- We have 5 parameters to worry about in convolution
 - `in_channels, out_channels`: How will the feature/channel dimension change?
 - `kernel_size`: Kernel is shape (k, k) typically
 - `padding`: How to deal with the edges? We almost always use “same” padding (or $k//2$) to have a better understanding of how shapes of tensors will change through the network
 - `stride`: Do we downsample by factor of s ?

Shapes in Convolution

- Channels
 - $(B, C_{in}, H, W) \rightarrow (B, C_{out}, H, W)$
- Padding, kernel size
 - If we use "same", then H and W don't change
- Stride
 - $(B, C, H, W) \rightarrow (B, C, H//s, W//s)$
 - Keyword "same" doesn't work, must use `padding=k//2` to ensure shapes change as expected
 - We typically downsample at certain points in the network, not often a hyperparameter

Regularization Activity

- Given this network that overfits, what should be done to improve its test performance?

Convolution Shapes Activity

- Work through how convolution affects input/output shapes (time permitting)